

# Society of Meta-Agents: The SOUL Motivation Framework for Specific and Objective Understanding Logic

Keyvan M. Sadeghi

April 22, 2025

## 1 Literature Review

We have identified six practical schools of agent motivation:

- **Intrinsic Motivation in RL:** Curiosity and surprise as auxiliary rewards to drive exploration in sparse environments (Schmidhuber, 2010; Pathak et al., 2017).
- **Information-Theoretic Drives:** Maximizing mutual information, novelty, and empowerment signals (Klyubin et al., 2005; Salge et al., 2014; Goertzel, 2024).
- **Competence-Progress Models:** Rewarding measurable learning progress toward self-generated subgoals to form a self-curriculum (Oudeyer & Kaplan, 2007).
- **Homeostatic/Drive-Reduction:** Minimizing internal variables (e.g., prediction error) via cybernetic drives (Hull, 1943; Friston, 2010).
- **Cognitive Architectures:** Embedding scalar drives into symbolic cycles (Soar’s operator preferences; ACT-R’s Expected Value of Control; Goertzel, 2024) for precise control updates.
- **Developmental-Robotics Hybrids:** Combining maturational constraints with competence progress to emulate developmental curricula (Lungarella et al., 2003).

## 2 Mathematical Foundations and Architectural Decisions

### 2.1 The Motivation Vector

At the heart of the SOUL Motivation Framework is the hidden internal state, the Motivation Vector:

$$\mathbf{s}_t = [s_c, s_u, s_h]$$

where  $s_c$  is competence,  $s_u$  is novelty/surprise, and  $s_h$  is homeostasis. This vector is updated after every interaction and drives all agent actions. In code, it is maintained as a Python dataclass and is never exposed to the LLM.

## 2.2 Motivation Vector Updates

- **Competence Progress:**

$$\Delta_c = p_g(t) - p_g(t-1), \quad s_c \leftarrow \text{clip}(s_c + \alpha \Delta_c, 0, 1)$$

where  $p_g(t)$  is the agent’s measured performance at time  $t$ .

- **Novelty/Surprise:**

$$\text{novel}(t) = 1 - \frac{\mathbf{e}(t) \cdot \mu_{t-1}}{\|\mathbf{e}(t)\| \|\mu_{t-1}\|}, \quad s_u \leftarrow \text{clip}(s_u + \alpha \text{novel}(t), 0, 1)$$

where  $\mathbf{e}(t)$  is the current context embedding and  $\mu_{t-1}$  is the rolling mean.

- **Homeostatic Decay:**

$$s_h \leftarrow (1 - \delta)s_h + \delta$$

This ensures  $s_h$  gently returns to baseline.

## 2.3 Meta-Graph and Rule Engine

The agent maintains a symbolic meta-graph  $G$  (a directed graph of rewrite rules  $R$  and meta-rules  $M$ ). Each node represents a pattern or rule, and edges encode transformations or relationships. In Python, this is implemented with `networkx.DiGraph`.

## 2.4 Thresholded Nudge and Confidence

At each step, the agent computes a confidence score  $C_t$ :

$$C_t = f_{\text{match}}(x_t, G, \mathbf{s}_t)$$

where  $f_{\text{match}}$  is a similarity or density function over meta-graph patterns and the current state. The agent compares  $C_t$  to a dynamic threshold  $\tau_t$ :

If  $C_t \geq \tau_t$  then nudge; else remain silent (null action)

**Concrete Example:** If  $C_t$  is computed as the softmax of pattern matches in  $G$  weighted by  $\mathbf{s}_t$ , and  $\tau_t$  is set adaptively (e.g., as a running quantile), then the agent only acts when it is sufficiently confident.

## 2.5 Perception–Cognition–Action Loop

The agent’s operation at each time  $t$  is:

$$\begin{aligned} \text{Perceive: } & y_t = \text{Perceive}(x_t) \\ \text{Update: } & \mathbf{s}_{t+1} = \text{UpdateMotivation}(\mathbf{s}_t, y_t) \\ \text{Record: } & G_{t+1} = \text{UpdateMetaGraph}(G_t, x_t, y_t) \\ \text{Confidence: } & C_{t+1} = f_{\text{match}}(x_{t+1}, G_{t+1}, \mathbf{s}_{t+1}) \\ \text{Action: } & a_{t+1} = \begin{cases} \text{HarvestAxiom}(G_{t+1}, x_{t+1}) & \text{if } C_{t+1} \geq \tau_{t+1} \\ \emptyset & \text{otherwise} \end{cases} \end{aligned}$$

*Explanation:*  $a_{t+1}$  is either a harvested axiom/rule (to be injected as a nudge) or the null action  $\emptyset$  (agent remains silent).

## 2.6 Subgoal Discovery

If  $s_c$  stagnates or  $s_u$  spikes, the agent auto-discovers new subgoals by clustering novel contexts in  $G$  and generating new rules. This enables adaptive exploration.

## 2.7 Discrete Generative Core and Error Signals

Instincts and policies are encoded as rewrite rules in  $G$ . The agent predicts a distribution  $p_t(m)$  over outcomes, observes  $q_t(m)$ , and computes error:

$$e_t = D_{\text{KL}}(q_t \| p_t)$$

This error drives reward and learning:

$$r_t^{\text{int}} = -e_t, \quad r_t^{\text{ep}} \propto \sum_m q_t(m) \log \frac{1}{p_t(m)}$$

## 2.8 Meta-Rule Self-Modification

Meta-rules  $M$  rewrite the rule graph itself:

$$m_i \leftarrow \arg \min_{m' \in \mathcal{N}(m_i)} e_t(R, M \setminus \{m_i\} \cup \{m'\})$$

## 2.9 Wasserstein Natural Gradient

Parameterize rule-distribution  $p(\xi)$  and update via:

$$\xi_{k+1} = \xi_k - h G(\xi_k)^{-1} \nabla_{\xi} F(p(\xi_k))$$

where  $G$  is the Laplacian from the rule graph.

## 2.10 Neural-Symbolic Hybrid and Memory

Continuous predictive-coding nets (vision and motor) run beneath the discrete core, exchanging features/actions. Long-term memory is implemented via vector stores (e.g., **faiss**, **chromadb**) for retrieval and adaptation.

## 2.11 LLM Pre-Prompting and Naturalization

When the agent nudges, it injects symbolic axioms/rules into the LLM prompt. The LLM is pre-prompted to interpret these in MeTTa or similar syntax and translate their intent into natural language or actions.

## 2.12 Genetic Mixing and Policy Sharing

Hyperparameters  $(\alpha, \delta, \tau_c, \tau_u)$  are encoded as arrays and can be evolved via genetic algorithms (e.g., DEAP, pygad), enabling agent societies to mix and share policies and meta-graphs.

## 2.13 Concrete Python Mapping

- **Motivation Vector:** Python `dataclass` with fields for  $s_c, s_u, s_h$ .
- **Rule Graph:** `networkx.DiGraph` with nodes for rules/meta-rules.
- **Neural Nets:** `torch.nn.Module` or `jax` models for predictive coding.
- **Memory:** `faiss` or `chromadb` vector store.
- **Hyperparameters:** Numpy array or genetic algorithm chromosome.

## 3 Null Action and Silent Learning

*If the confidence  $C_t$  does not exceed the threshold  $\tau_t$ , the agent performs the null action  $\emptyset$ , i.e., it remains silent and continues to observe, record, and learn without intervening.*

## 4 Implementation Guide

Below are concrete steps, with Python library suggestions.

- **Competence-Progress Core:** Track competence gains  $\Delta_c$  on self-generated sub-goals, updating  $s_c \in [0, 1]$  via

$$s_c \leftarrow \text{clip}(s_c + \alpha \Delta_c, 0, 1)$$

- **Novelty/Surprise Seeding:** Compute novelty as cosine distance of new embedding  $\mathbf{e}(t)$  to recent mean  $\mu_{t-1}$ ,

$$\text{novel}(t) = 1 - \frac{\mathbf{e}(t) \cdot \mu_{t-1}}{\|\mathbf{e}(t)\| \|\mu_{t-1}\|}, \quad s_u \leftarrow \text{clip}(s_u + \alpha \text{novel}(t), 0, 1)$$

- **Homeostatic Decay:** Maintain stability  $s_h$  toward 1 via

$$s_h \leftarrow (1 - \delta) s_h + \delta$$

- **Thresholded Nudge Mechanism:** At each step, the agent computes a confidence score  $C_t$  based on the match between the current context  $x_t$ , the meta-graph  $G$ , and the motivation vector  $\mathbf{s}_t$ :

$$C_t = f_{\text{match}}(x_t, G, \mathbf{s}_t)$$

The agent compares  $C_t$  to a dynamic threshold  $\tau_t$ :

If  $C_t \geq \tau_t$  then nudge; else remain silent (null action)

*Explanation:*  $f_{\text{match}}$  may be a similarity or density function over meta-graph patterns, and  $\tau_t$  can be static or adaptively tuned.

- **Perception–Cognition–Action:** As formulated in Section 2.5, the agent computes and thresholds a confidence score to determine whether to nudge or remain silent. This is implemented by evaluating a match function between the current context, meta-graph, and motivation vector, and comparing it to a dynamic threshold.
- **Discrete Generative Core:** Represent “instincts” as a *metagraph* of rewrite rules  $R$ , inducing a predicted distribution  $p_t(m)$  over outcomes. Measure error [1]

$$e_t = D_{\text{KL}}(q_t \| p_t)$$

- **Reward Signals:** Combine instrumental  $r_t^{\text{int}} = -e_t$  and epistemic  $r_t^{\text{ep}} \propto \sum_m q_t(m) \log \frac{1}{p_t(m)}$  to guide local rule edits.[1]
- **Meta-Rule Self-Modification:** Define meta-rules  $M$  that pattern-match on  $R$  and refactor complex rules. Local meta-update:

$$m_i \leftarrow \arg \min_{m' \in \mathcal{N}(m_i)} e_t(R, M \setminus \{m_i\} \cup \{m'\}) \quad [1]$$

Meta-rules undergo the same reward-driven selection as object-level rules.

- **Wasserstein Natural Gradient:** Parameterize rule-distribution  $p(\xi)$  and update via the optimal-transport natural gradient:[1]

$$\xi_{k+1} = \xi_k - h G(\xi_k)^{-1} \nabla_{\xi} F(p(\xi_k))$$

with metric tensor  $G$  from the measure-dependent Laplacian on the rule graph.

- **Neural–Symbolic Hybrid:** Two continuous predictive-coding nets (vision and motor) run beneath the discrete core (a *metagraph* of self-transforming codelets), passing symbolic features and actions in a closed-loop.[1]

**Null Action and Silence:** *If the confidence  $C_t$  does not exceed the threshold  $\tau_t$ , the agent performs the null action  $\emptyset$ , i.e., it remains silent and continues to observe, record, and learn without intervening.*

## 5 Implementation Guide

Below are concrete steps, with Python library suggestions.

### 5.1 Core Data Structures

- Use `numpy` for vector operations and embeddings.
- Use `networkx` to represent the metagraph of rewrite rules and compute Laplacians.

- Store agent state in a simple `dataclass`:

```
from dataclasses import dataclass
@dataclass
class State:
    competence: float = 0.0
    curiosity: float = 0.0
    stability: float = 1.0
```

## 5.2 Novelty Detector

- Implement rolling mean with `collections.deque` and cosine distance via `scipy.spatial.distance.cos`

## 5.3 Rewrite-Rule Engine

- Model rules as Python objects mapping pattern graphs to outputs.
- Use `networkx` pattern-matching or custom graph algorithms for rule application.
- Derive  $p_t(m)$  by sampling or counting rule firings over stochastic selections (e.g., softmax weights in `torch`).

## 5.4 Information-Theoretic Error

- Compute KL divergence with `scipy.stats.entropy(q, p)`.

## 5.5 Reward & State Update Loop

1. Collect feedback score  $r(t)$  via environment simulation or user rating.
2. Update `State` (Motivation Vector) using the mathematical formulas above, with learning rate `alpha`.
3. Compute confidence  $C_t$  and compare to threshold  $\tau_t$ ; if  $C_t \geq \tau_t$ , harvest and inject a nudge (axiom/rule) from the meta-graph.
4. If nudging, prepend the harvested axiom/rule to the user query and invoke the LLM (e.g., via `openai` or `transformers`), relying on a pre-prompt for interpretation.
5. If not nudging, remain silent and continue to observe, record, and update internal state.

**Example:** In Python, this loop is implemented as a function that updates the Motivation Vector, computes confidence, and either calls a nudge-injection routine or skips to the next observation.

## 5.6 Meta-Rule Implementation

- Represent meta-rules as rules over the rule-graph using `networkx`.
- Define neighborhood  $\mathcal{N}(m_i)$  of small metagraph edits.
- Apply the meta-rule update formula in your training loop alongside rule edits.

## 5.7 Natural Gradient Optimization

- Install `pot` (Python Optimal Transport) for Wasserstein solvers.
- Build ground metric matrix  $(\omega_{ij})$  from rule-graph distances.
- Construct measure-dependent Laplacian via NetworkX weights.
- Compute parameter Jacobians with `autograd` or manual derivatives.
- Perform updates using `numpy.linalg.pinv` for pseudoinverse.

## 5.8 Continuous Predictive-Coding Nets

- Use PyTorch or JAX to implement NGC-style layers:

$$z \leftarrow z + \beta(-\gamma z + (E \cdot e) \odot \phi'(z) - e)$$

- Leverage `torch.nn.Module` for vision and arm nets.
- Optimize with local Hebbian-like or standard optimizers (SGD) per PC update.

## 5.9 Genetic Tuning (Optional)

- Represent hyperparameters  $(\alpha, \delta, \tau_c, \tau_u)$  as a NumPy array.
- Use DEAP or pygad for crossover and mutation over logged performance metrics.

## 5.10 Vector Stores and Long-Term Memory

- Integrate `faiss` or `chromadb` for storing past contexts/embeddings.

# 6 Conclusion

This guide unifies seminal research, architectural principles, and practical Python tools to implement SOUL’s Motivation Framework. By following it, you’ll build agents that self-modify, learn from surprise and progress, and nudge LLMs with precise, adaptive prompts.

## References

- [1] B. Goertzel. Discrete Active Predictive Coding for Goal-Guided Algorithmic Chemistry as a Potential Cognitive Kernel, 2024. arXiv:2412.16547 [cs.AI].
- [2] D. Pathak *et al.*, Curiosity-Driven Exploration by Self-Supervised Prediction, 2017.
- [3] C. Salge *et al.*, Curiosity-Driven RL Using Information-Gain, 2014.
- [4] K. Friston. The Free-Energy Principle: A Unified Brain Theory?, 2010.
- [5] J. Schmidhuber. Formal Theory of Creativity, Fun, and Intrinsic Motivation, 2010.
- [6] P.-Y. Oudeyer, F. Kaplan, and V. V. Hafner. Intrinsic Motivation Systems for Autonomous Mental Development, 2007.
- [7] A. S. Klyubin *et al.*, Empowerment: A Universal Agent-Centric Measure, 2005.
- [8] M. Lungarella, G. Metta, R. Pfeifer, and G. Sandini. Developmental robotics: A survey, *Connection Science*, 15(4), 151-190, 2003.
- [9] C. L. Hull. *Principles of Behavior: An Introduction to Behavior Theory*. New York: Appleton-Century, 1943.